

Using Terraform to Deploy Hadoop Infrastructure

Presented By Michael Arnold

CLAIRVOYANT

Agenda

- About the speaker
- What
- How

Who is Michael Arnold?

- Principal Systems Engineer/Consultant
- Automation geek
- 20 years in IT - 8 of those years with Hadoop
- I help people deal with:
 - Servers (physical and virtual)
 - Networks
 - Server operating systems
 - Hadoop distributions
 - Making it all run smoothly

What is Terraform?

- A tool for building infrastructure as code.
- A declarative language.
- Allows for visualizing changes that will be made.
- Automatically determines resource dependencies.
- Can build servers/networks/databases in the cloud.
- Can change DNS records, create git repos, configure monitoring services, or generate TLS certificates.

Resources

- Providers
 - Services which provide an API to which Terraform can communicate.

Providers



Resources

- Provisioners
 - Scripts or commands that can be run to bootstrap a resource.

Provisioners

connection
salt_masterless
remote_exec
local_exec
chef
file
null_resource

Commands

- `terraform init` # only needed once
- `terraform plan`
- `terraform apply`
- `terraform destroy`

Demonstration

DEMO Time!

Other Commands

- `terraform graph`

Code

- <https://github.com/teamclairvoyant/terraform-hadoop-talk>
- <https://github.com/teamclairvoyant/hadoop-deployment-bash>
- <https://www.terraform.io/>

Michael Arnold
@hadoopgeek
www.linkedin.com/in/michaelarnold

www.clairvoyantsoft.com