

# Airflow

## Agenda

- Introduction to Airflow
- Architectural Overview
- Deployment Models
- Setting up Airflow for High Availability
- Common issues
- Demo of important features
- Airflow in production - a few use cases
- QA

# Introduction

## **Data engineering workflow management - requirements/challenges**

Connect to Data Sources

Orchestration of various stages in the pipeline dynamically with no downtime

Managing dependencies between stages

Scheduling jobs

Monitoring the health of the pipelines

Hooks in Airflow

Pipelines are configured via code making the pipelines dynamic

XCom

Has support for Calendar schedule and Cron tab scheduling

A graphical representation of the DAG instances and Task Instances along with the metrics.  
Email notifications are natively supported

# Introduction

## **Data engineering workflow management - requirements/challenges**

Fault Tolerant pipelines

Retry Mechanism and Backfill utilities in Airflow

Controlling the Pipeline execution externally

PlugIn architecture that helps to add new functionality (REST plugin), TriggerDags

Scaling the Workflow Management System

Scalability: Distribution of Workers and Queues for Task execution

More control to the Management/Analysts to control the pipelines

Variables for making the changes to the Dags/Tasks quick and easy

Security

kerberos Support

# Getting Started

- DAGs —> Collection/Orchestrated Tasks
- Executors
  - derives from BaseExecutor
  - Three types: Sequential, Local and Celery
- Schedule properties
  - start\_date
  - end\_date
  - schedule\_interval
  - depends\_on\_past

# Getting Started

## → Operators

- derives from BaseOperator
- Three types: Sensors, RemoteExecution, Data Transfer
  - ExternalTaskSensor
- trigger\_rule --> defines when the task has to be triggered based on the status of upstream tasks.

## → Hooks

- derives from BaseHook.
- Operators use hooks that actually has API calls to perform an operation

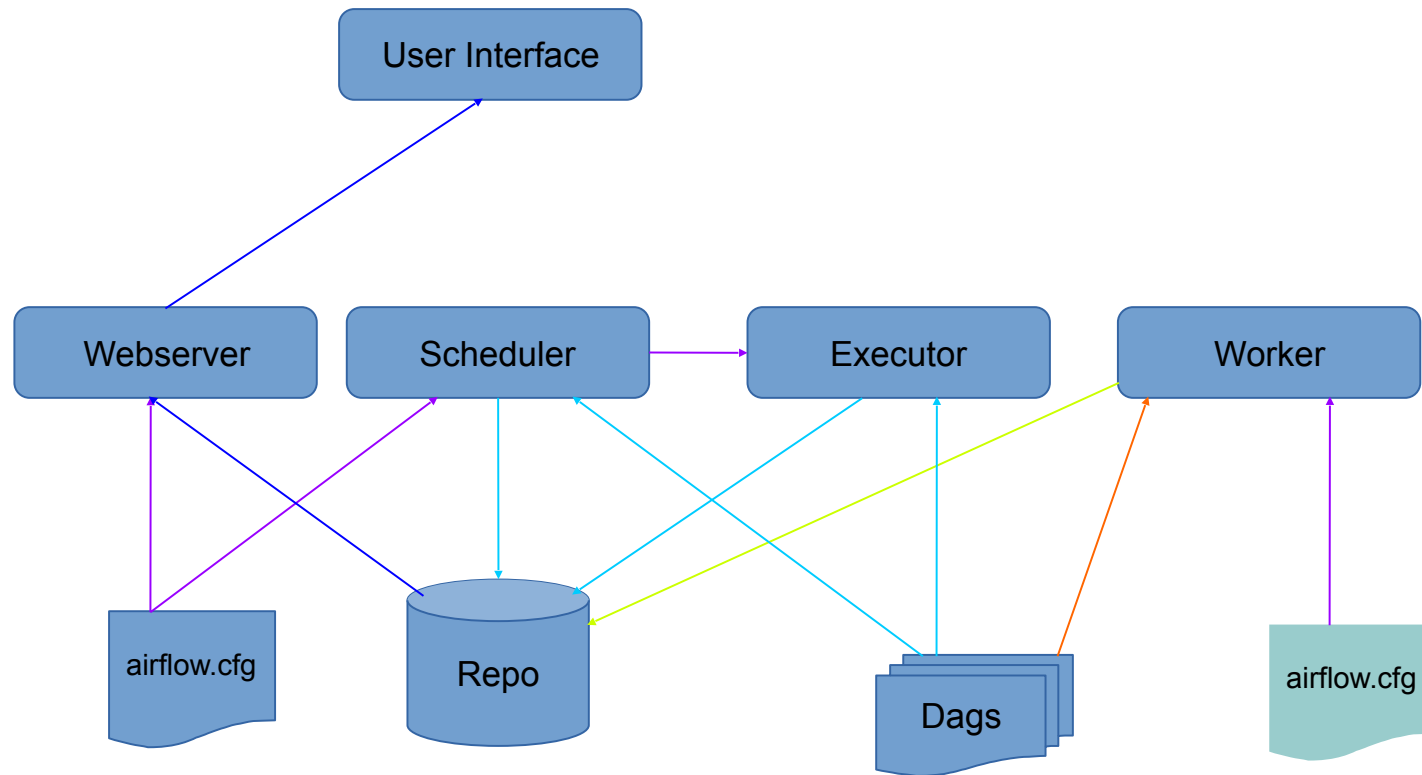
## → Task

- a parameterized Operator
- Limiting Parallelism: Supports LIMIT on parallelism and PRIORITY\_WEIGHT for tasks.
- Dependencies – upstream()/downstream()

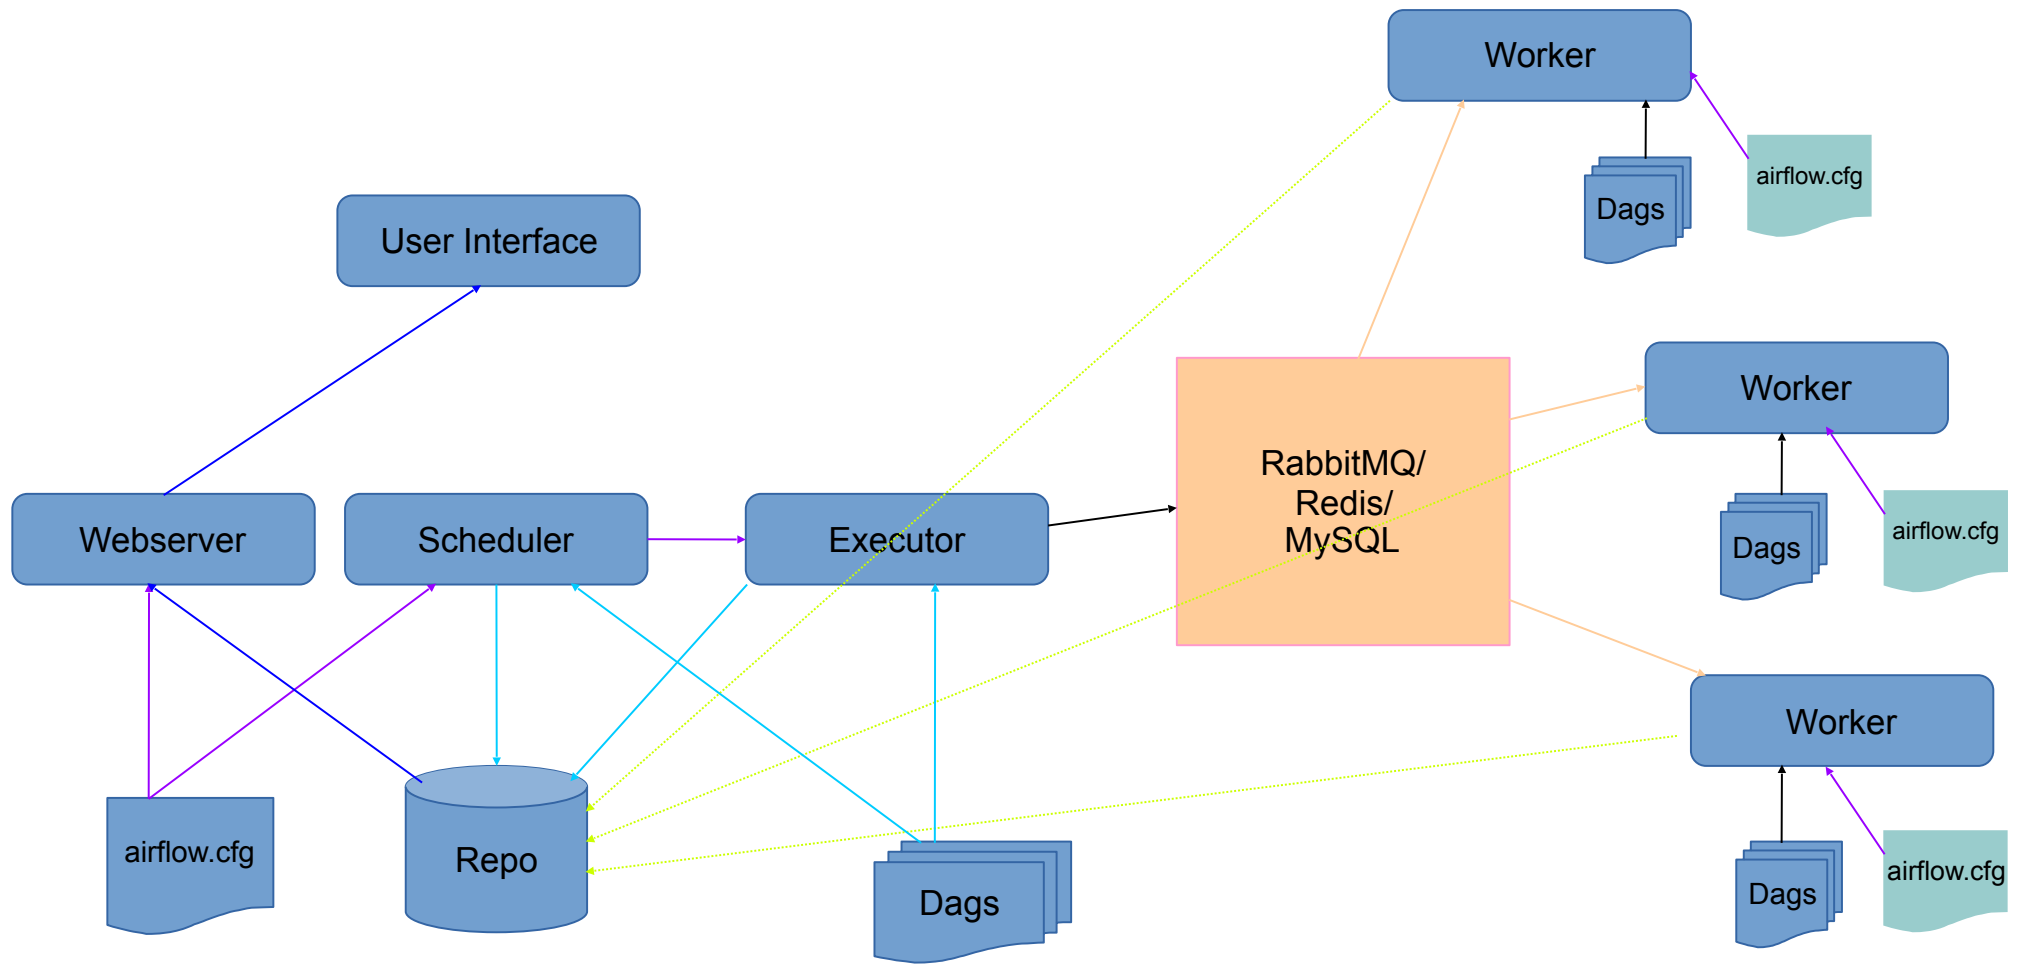
# Components of Airflow

- Configuration file
- a metadata database (mysql or postgres)
- the Airflow scheduler
- a broker (redis or rabbitmq)
- a set of Airflow worker nodes
- the Airflow web server

# Components of Airflow

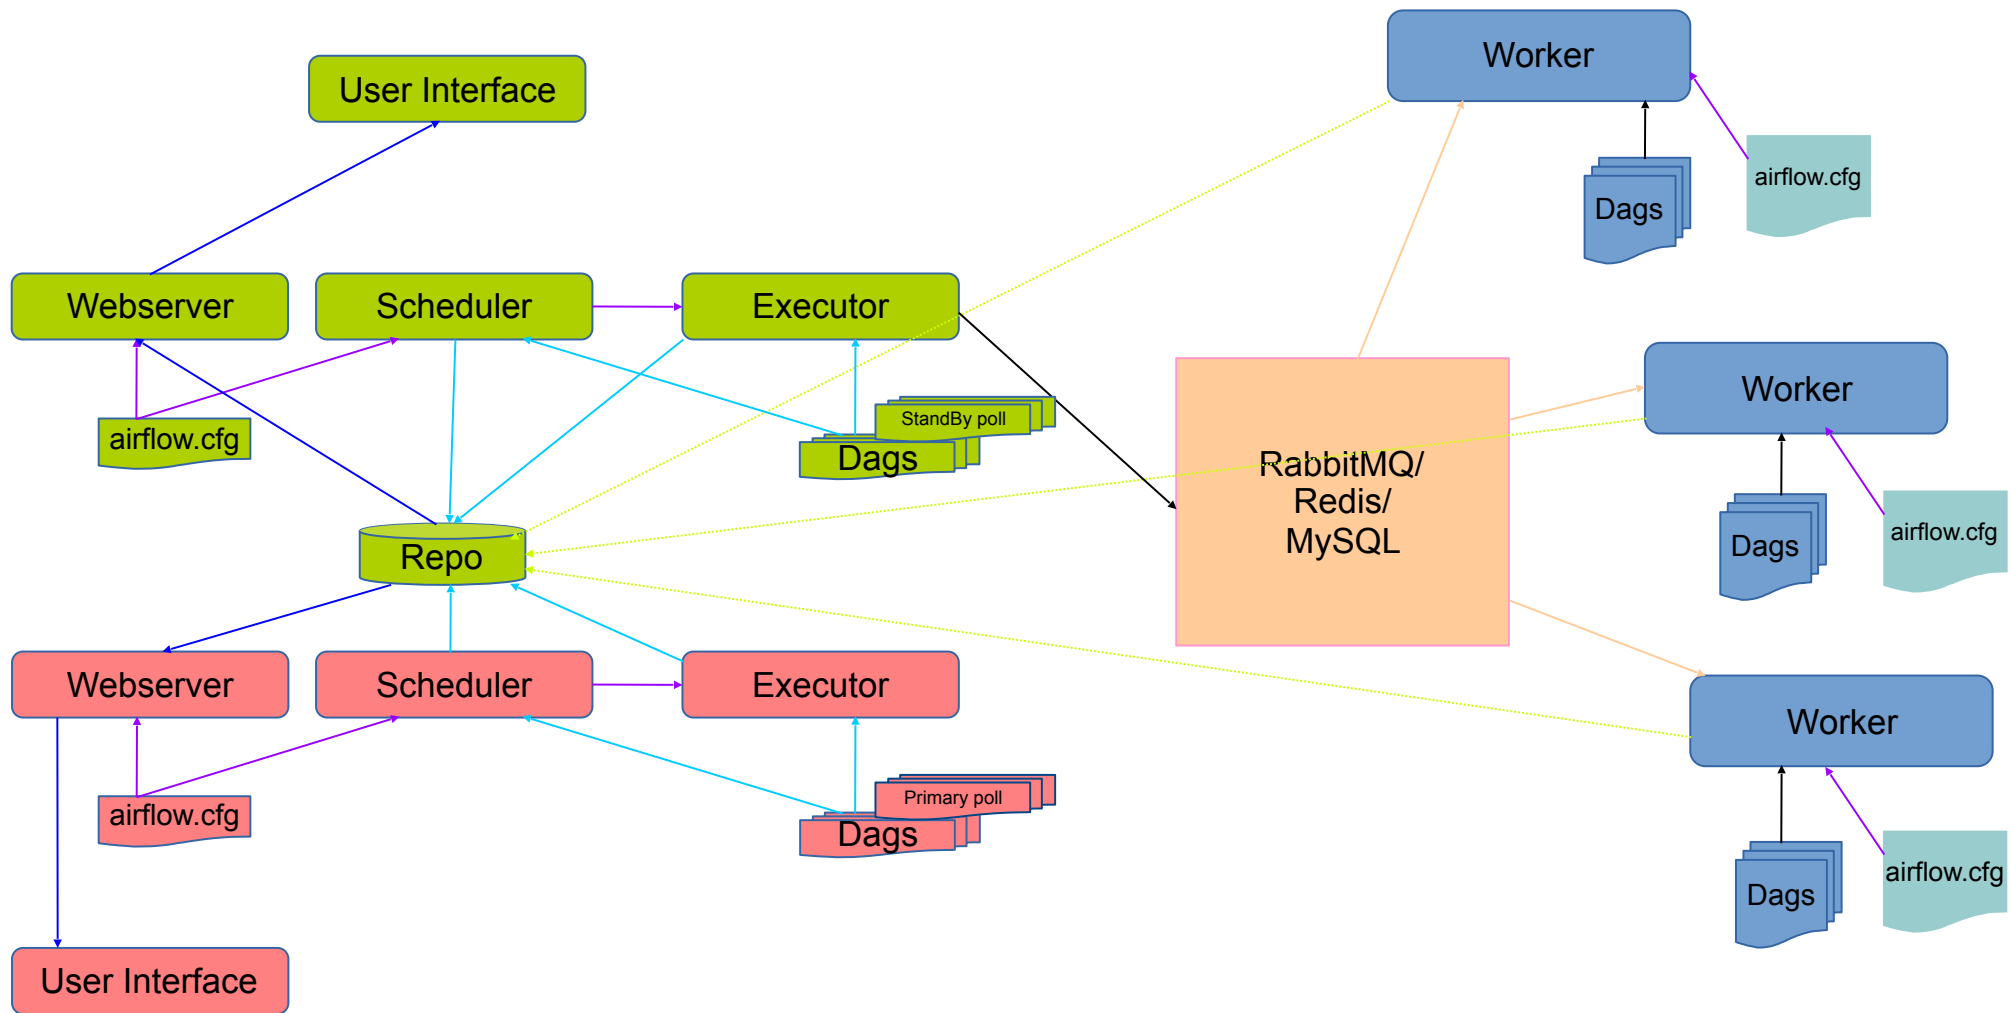


# Components of Airflow

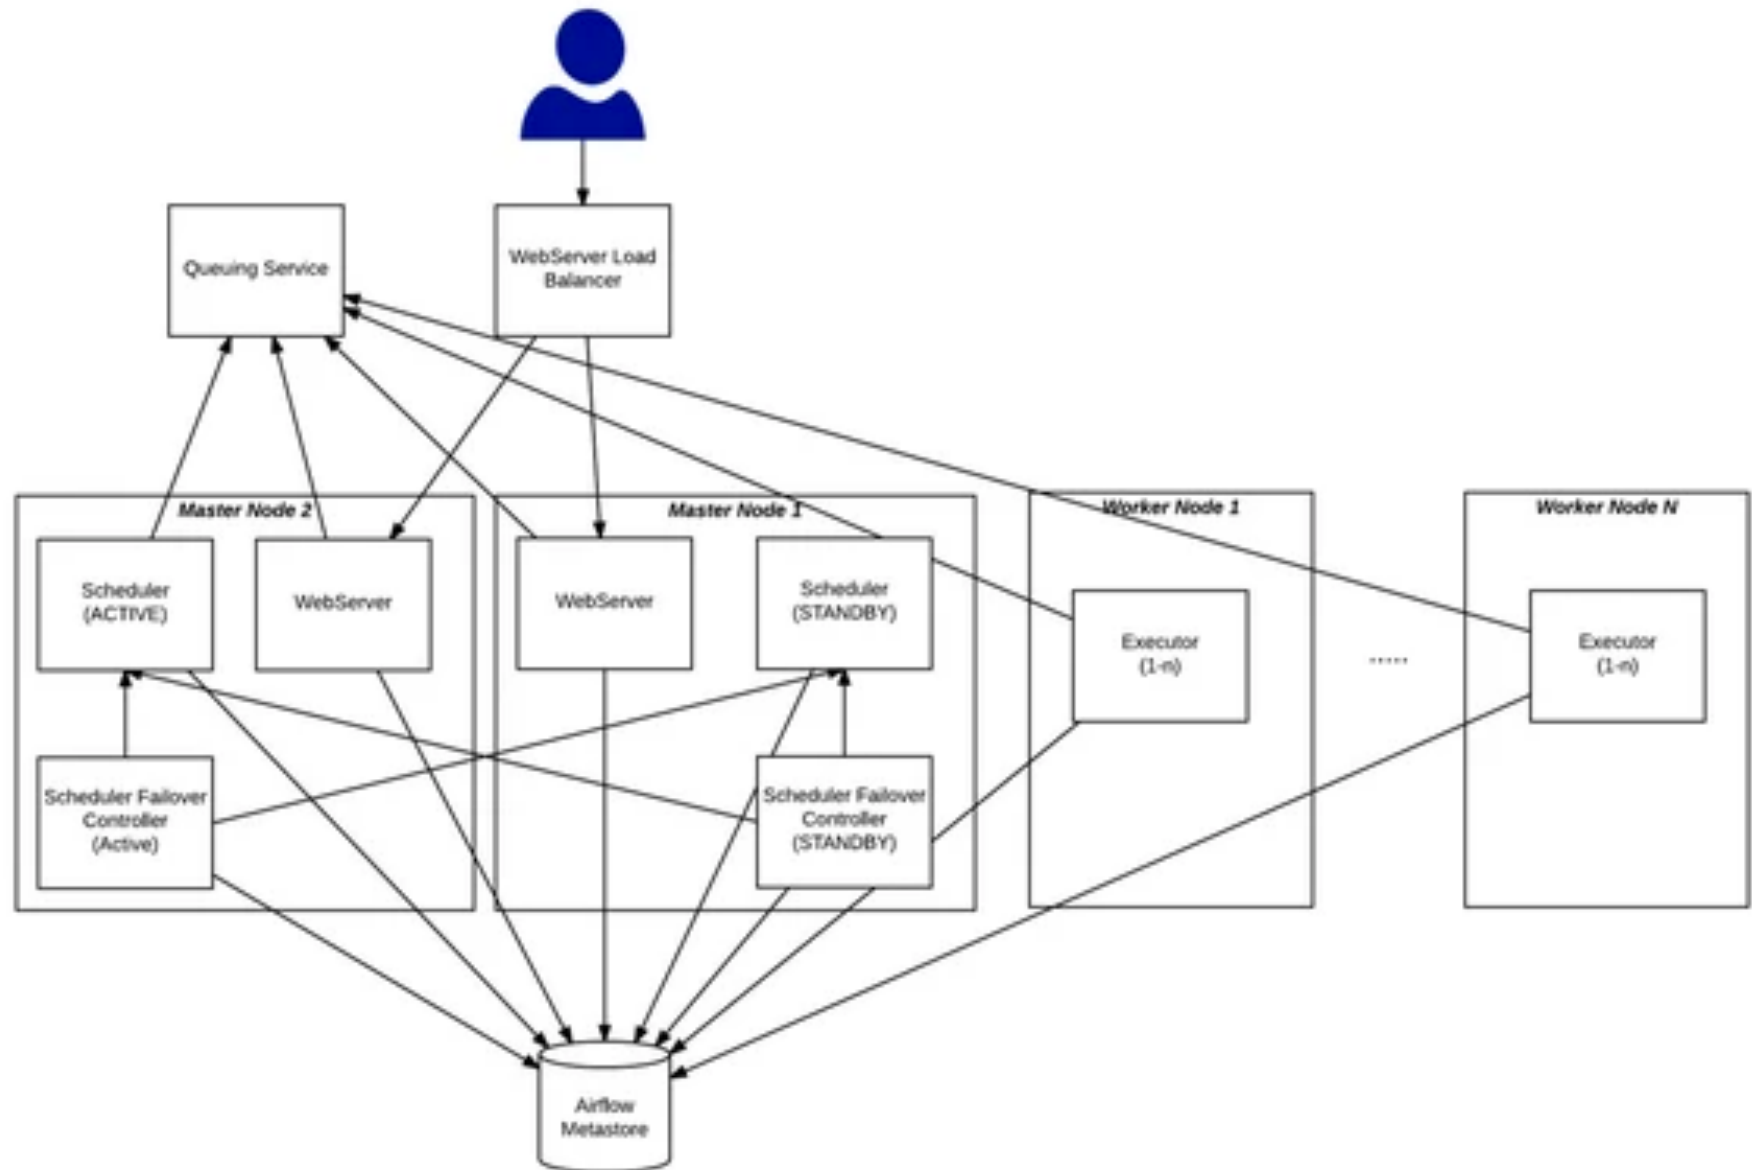




# Components of Airflow



# Components of Airflow



# Common Issues

1. Unintended instances when the DAG is not executed for a couple of days.
2. Limit on number of Tasks in a DAG - UI issues
3. If DB goes down, no clue to the end user.
4. Historic instances refer to latest DAG definition

# Important points

- RabbitMQ is to be configured for Celery workers to connect to the RabbitMQ server.
- Airflow/Celery needs to be installed on every server.
- The codebase for the tasks and the scripts that are being called from Airflow needs to be present on all the workers.
  - The code based should be in sync to avoid any un-intended results.
- RabbitMQ can be run in standalone mode or cluster mode. Cluster mode would ensure that the queues are distributed.

# Demo

# Some links

Steps to install RabbitMQ on Centos

<http://geek-kb.com/rabbitmq-server-centos-6x/>

HA for Rabbit MQ

<https://www.rabbitmq.com/blog/2011/10/25/high-availability-in-rabbitmq-solving-part-of-the-puzzle/>

Installing Mysql in Centos

<https://www.linode.com/docs/databases/mysql/how-to-install-mysql-on-centos-7>

Replicating MySQL Database

<https://www.digitalocean.com/community/tutorials/how-to-set-up-master-slave-replication-in-mysql>